

Text Analytics

Ramaseshan Ramachandran

January 25, 2025

WHAT IS TEXT ANALYTICS?



- ▶ Converts unstructured text into insights
- ▶ Extracts patterns and trends
- ▶ Uses NLP, machine learning, and statistics
- ▶ Analyzes data from diverse sources

- ▶ Unstructured text data is vast
- ▶ Text data is messy and variable
- ▶ Makes data measurable and valuable

- ▶ Understand customers and societal trends
- ▶ Enhance decision-making and planning
- ▶ Drive engagement and innovation

- ▶ Searches documents and metadata
 - ▶ Extracts relevant data to queries
 - ▶ Applies to large text collections
- ▶ Boolean Retrieval Model
 - ▶ Vector Space Model (TF-IDF)
 - ▶ Probabilistic Retrieval Model
 - ▶ Latent Semantic Indexing (LSI)
 - ▶ BM25 Algorithm

- ▶ Categorizes text into predefined labels
 - ▶ Applications: spam detection, sentiment analysis
 - ▶ Automates document organization
- ▶ Naive Bayes Classifier
 - ▶ Support Vector Machines (SVM)
 - ▶ Logistic Regression
 - ▶ Decision Trees and Random Forests
 - ▶ Deep Learning (CNNs, RNNs, Transformers)

- ▶ Groups similar texts together
 - ▶ No predefined labels required
 - ▶ Useful for exploratory analysis
- ▶ K-Means Clustering
 - ▶ Hierarchical Clustering
 - ▶ DBSCAN (Density-Based Clustering)
 - ▶ Gaussian Mixture Models (GMMs)
 - ▶ Spectral Clustering

- ▶ Determines text's emotional tone
 - ▶ Classifies as positive, negative, or neutral
 - ▶ Applications: marketing, feedback analysis
- ▶ Lexicon-Based Approaches
 - ▶ Rule-Based Sentiment Analysis
 - ▶ Machine Learning-Based Approaches
 - ▶ Neural Networks (LSTMs, GRUs, BERT)
 - ▶ Pretrained Models (RoBERTa, GPT)

- ▶ Identifies entities in text (e.g., names)
 - ▶ Categorizes into people, places, etc.
 - ▶ Useful for automated content analysis
- ▶ Hidden Markov Models (HMMs)
 - ▶ Conditional Random Fields (CRFs)
 - ▶ Maximum Entropy Models
 - ▶ Neural Networks (BiLSTM + CRF)
 - ▶ Pretrained Models (SpaCy, Hugging Face)

- ▶ Discovers themes in document collections
- ▶ Applications: summarization, recommendations

Algorithms

- ▶ Latent Dirichlet Allocation (LDA)
- ▶ Non-Negative Matrix Factorization (NMF)
- ▶ Latent Semantic Analysis (LSA)
- ▶ Gibbs Sampling for LDA
- ▶ Neural Topic Models (ProdLDA, BERTopic)

- ▶ Word Embedding Models (Word2Vec, GloVe, FastText)
- ▶ Sentence Embeddings (Sentence-BERT)
- ▶ Attention Mechanisms
- ▶ Transformers (BERT, GPT, T5)
- ▶ Text Summarization (Extractive and Abstractive)

- ▶ Text analytics unlocks data potential
- ▶ Drives decisions and innovation
- ▶ Benefits industries like finance and healthcare
- ▶ A variety of algorithms drive text analytics
- ▶ Techniques range from statistical to neural
- ▶ Tailor solutions based on use case

WORD EMBEDDING CANNOT FIGHT WITH OTHERS. WHY



- ▶ **Definition:** Scikit-learn is a Python library for machine learning, providing tools for:
 - ▶ Classification
 - ▶ Regression
 - ▶ Clustering
 - ▶ Dimensionality reduction
 - ▶ Preprocessing and more
- ▶ **Key Features:**
 - ▶ Built on NumPy, SciPy, and matplotlib
 - ▶ Simple and efficient tools for predictive data analysis
 - ▶ Open source

- ▶ **Text Analytics Focus:**

- ▶ Natural Language Processing (NLP) tasks
- ▶ Feature extraction (e.g., bag-of-words, TF-IDF)
- ▶ Building predictive models

- ▶ **Advantages:**

- ▶ Wide range of algorithms (SVMs, Naive Bayes, etc.)
- ▶ User-friendly API for rapid prototyping
- ▶ Extensive documentation and community support

▶ **Examples:**

- ▶ Spam detection
- ▶ Sentiment analysis
- ▶ Topic modeling
- ▶ Document classification

▶ **Techniques:**

- ▶ Preprocessing: Tokenization, stemming, lemmatization
- ▶ Vectorization: TF-IDF or CountVectorizer
- ▶ Model training: Logistic regression, SVMs, etc.

1. Data Preprocessing:

- ▶ Cleaning text data (e.g., removing stop words)
- ▶ Vectorization (e.g., TF-IDF)

2. Model Selection:

- ▶ Choosing an algorithm (e.g., Naive Bayes)

3. Model Training:

- ▶ `model.fit(X_train, y_train)`

4. Model Evaluation:

- ▶ Metrics like accuracy, precision, recall

5. Prediction:

- ▶ `model.predict(X_test)`

Dataset: Sentiment Analysis on Product Reviews

1. Load dataset
2. Preprocess text (lowercase, remove punctuation, etc.)
3. Vectorize with TF-IDF
4. Train model (e.g., Logistic Regression)
5. Evaluate using accuracy and F1-score

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

vectorizer = TfidfVectorizer()
X = vectorizer.fit_transform(text_data)
X_train, X_test, y_train, y_test = train_test_split(X, labels, test_size=0.2)
model = LogisticRegression()
model.fit(X_train, y_train)
predictions = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, predictions))
```

▶ **Strengths:**

- ▶ Easy to use and integrate
- ▶ Extensive support for text-related tasks
- ▶ Scalability for moderate-sized datasets

▶ **Limitations:**

- ▶ Not designed for deep learning
- ▶ Limited support for out-of-core learning

- ▶ Official Documentation: <https://scikit-learn.org>
- ▶ Tutorials:
 - ▶ "Getting Started with Scikit-Learn" (Blog/Video)
 - ▶ Kaggle courses on ML
- ▶ Recommended Books:
 - ▶ *Python Machine Learning* by Sebastian Raschka
 - ▶ *Introduction to Machine Learning with Python* by Andreas Müller

Questions?